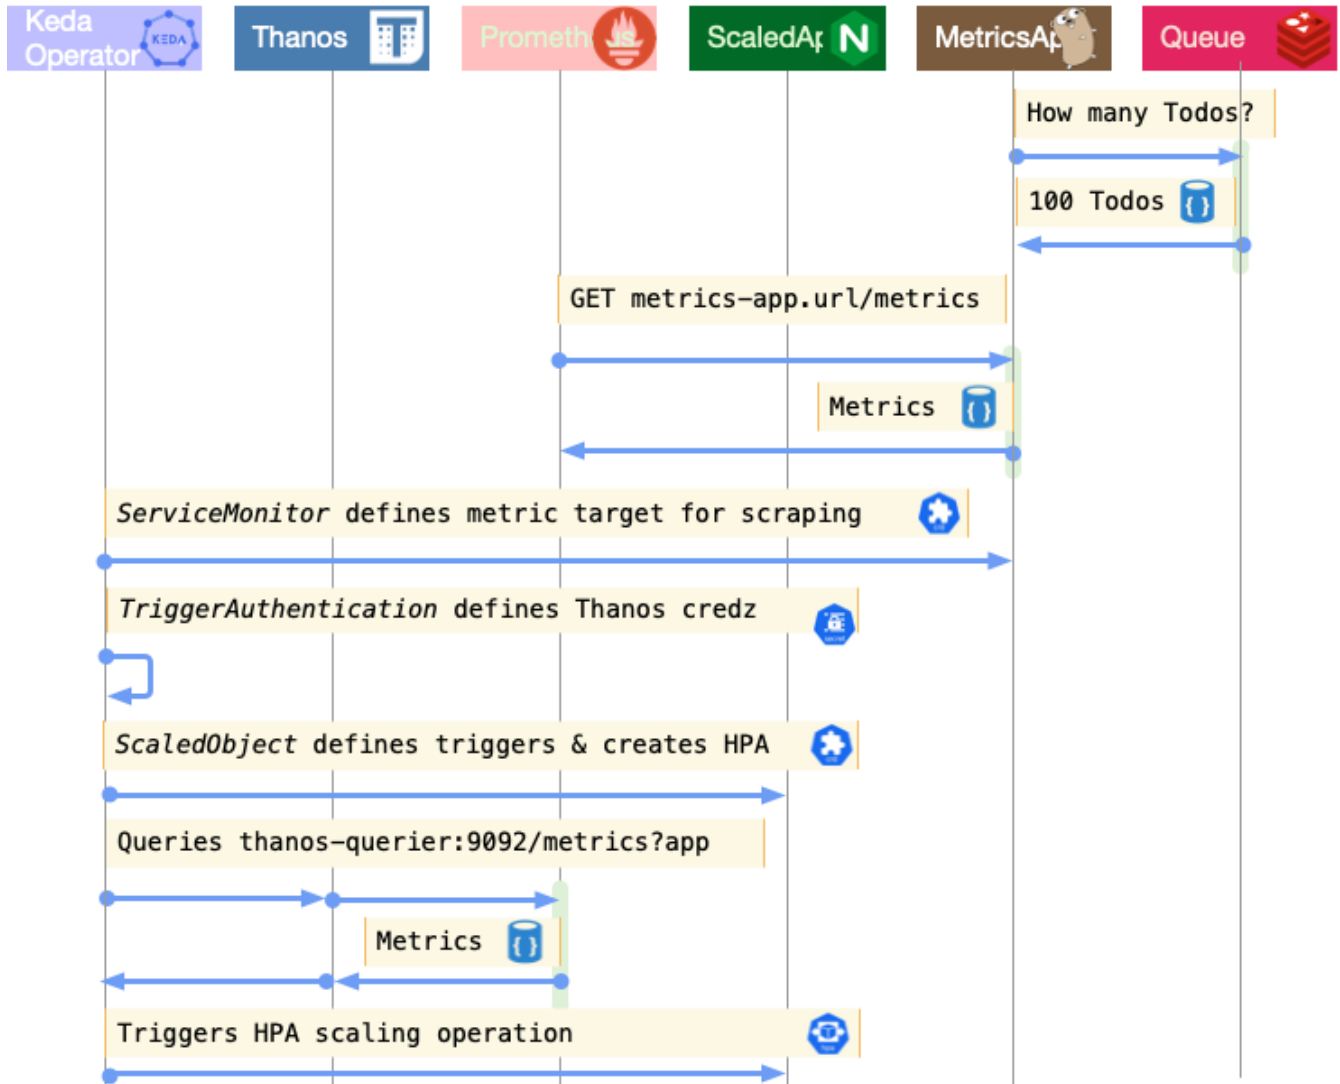


# Demo Using Custom Prometheus Metrics for OpenShift Autoscaling with KEDA



Updated: 2022-10-11

## Outline of Steps

- Enable Prometheus user workload monitoring
- Install the Custom Metrics Autoscaler Operator and provision a KedaController
- Deploy an application with custom metrics target and an associated ServiceMonitor
- Create Thanos serviceaccount, token, role, rolebinding
- Configure authentication from Keda to Thanos
- Deploy sample application to autoscale based on above metrics
- Define deployment to be scaled and the metrics which trigger the scaling with ScaledObject

## Docs and Refs

- <https://cloud.redhat.com/blog/custom-metrics-autoscaler-on-openshift>
- [Knative versus KEDA](#)
- [Enabling user workload monitoring in OpenShift](#)
- <https://docs.openshift.com/container-platform/4.11/nodes/pods/nodes-pods-autoscaling-custom.html>
- [Autoscaling for RGW in ODF via HPA using KEDA](#)
- <https://docs.openshift.com/container-platform/4.11/monitoring/managing-metrics.html>
- <https://github.com/rhobs/prometheus-example-app>
- <https://keda.sh/docs/2.8/scalers/prometheus/>
- [Enabling TLS in ServiceMonitor](#)
- <https://github.com/marrober/pipelineBuildExample>

# Prerequisites

## Enable User Workload Monitoring

- [OpenShift user workload monitoring](#)

```
oc extract configmap/cluster-monitoring-config \
  -n openshift-monitoring --to=- \
  | yq eval '.enableUserWorkload = true' - > config.yaml
oc set data configmap/cluster-monitoring-config \
  --from-file=config.yaml -n openshift-monitoring
```

## Confirm User Workload Monitoring

Check that pods for prometheus-operator, prometheus-user-workload and thanos-ruler-user-workload are running in the openshift-user-workload-monitoring project.

## Granting non-admin users permission to monitor user-defined projects (optional)

### Roles

- **monitoring-rules-view** grants read access to PrometheusRule custom resources for a project.
- **monitoring-rules-edit** grants create, modify, and deleting PrometheusRule custom resources for a project.
- **monitoring-edit** grants **monitoring-rules-edit** plus create new scrape targets for services or pods. With this role, you can also create, modify, and delete ServiceMonitor and PodMonitor resources.

Grant `monitoring-edit` role.

## Deploy Custom Metrics Autoscaler Operator

Deploys the OpenShift [Custom Metric Autoscaler operator](#) which is downstream from [KEDA.sh](#).

```
oc apply -k operator
namespace/openshift-keda created
kedacontroller.keda.sh/keda created
operatorgroup.operators.coreos.com/openshift-keda created
subscription.operators.coreos.com/openshift-custom-metrics-autoscaler-operator
created
oc logs -f -n openshift-keda -l app=keda-operator
```

## Deploy Custom Metric Producer

This application exists to expose a metrics endpoint that proffers custom metrics used for scaling a [second application](#).

This metric may be tangentially related to the load of the scaled app. For instance it may gather and report on a queue depth or work backlog. However, the current [app referenced here](#) does not.

```
oc apply -k custom-metric-app
namespace/keda-test created
service/prometheus-example-app created
deployment.apps/prometheus-example-app created
servicemonitor.monitoring.coreos.com/prometheus-example-monitor created
route.route.openshift.io/prometheus-example-app created
```

## View Example metrics

```
ROUTE=$(oc get route prometheus-example-app -o jsonpath='{.spec.host}')
curl $ROUTE/metrics
# HELP version Version information about this binary
# TYPE version gauge
version{version="v0.4.1"} 1
# create some traffic stats then retry above
curl $ROUTE
```

## TODO

### TIP

Create a new app with more pertinent metric example. Like returning the number of Todos from a work queue.

# Deploy Scaled-App

This is a sample application that is scaled by virtue of custom metrics exported from [custom metric app](#).

```
oc apply -k scaled-app
namespace/keda-test unchanged
serviceaccount/thanos created
role.rbac.authorization.k8s.io/thanos-metrics-reader created
rolebinding.rbac.authorization.k8s.io/thanos-metrics-reader created
secret/thanos-token created
service/static-app created
deployment.apps/static-app created
buildconfig.build.openshift.io/static-app created
imagestream.image.openshift.io/static-app created
scaledobject.keda.sh/static-app created
triggerauthentication.keda.sh/keda-trigger-auth-prometheus created
route.route.openshift.io/static-app created
```

```
oc get deployment,scaledobject,hpa -n keda-test
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/prometheus-example-app	1/1	1	1	37m
deployment.apps/static-app	1/1	1	1	36m

NAME	SCALETARGETKIND	SCALETARGETNAME	MIN	MAX
TRIGGERS AUTHENTICATION	READY ACTIVE FALLBACK	AGE		
scaledobject.keda.sh/static-app	apps/v1.Deployment	static-app	1	10
prometheus	keda-trigger-auth-prometheus	True False False	36m	

NAME	REFERENCE
TARGETS MINPODS MAXPODS REPLICAS AGE	
horizontalpodautoscaler.autoscaling/keda-hpa-static-app	Deployment/static-app 0/5
(avg) 1 10 1 36m	

## Managing metrics targets

As admin Observe → Targets → filter to 'keda-test' namespace.

## Metrics Targets

Filter ▾

Text ▾ keda-test

Text keda-test X

Clear all filters

Endpoint ▴	Monitor	Status ▴	Namespace ▴	Last Scrape ▴	Scrape Dura... ▴
<a href="http://10.128.5.174:8080/metrics">http://10.128.5.174:8080/metrics</a>	prometheus-example-monitor	Up	keda-test	Just now	3.3 ms

Target <http://<ip-address>:8080/metrics> corresponds to endpoint of service 'prometheus-example-app'

Go to Observe → Metrics → promql query and search for "version" to see metric from [custom metric app](#).

version

X

Name ▴	endpoint ▴	instance ▴	job ▴	namespace ▴	pod ▴	prometheus ▴	service ▴	version ▴	Value ▴
version	web	10.128.5.174:8080	prometheus-example-app	keda-test	prometheus-example-app-6dcf7dcf8b-hzb19	openshift-user-workload-monitoring/user-workload	prometheus-example-app	v0.4.1	1

<https://docs.openshift.com/container-platform/4.11/monitoring/managing-metrics-targets.html>

# Testing Autoscaling

Create a job that will generate traffic to the [custom metric app](#). When the rate of traffic hits a 1 minute rolling average of 5 hits then the autoscaler will scale out the [static-app deployment](#).

```
oc create -f load.yaml
# wait roughly a minute...
oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
generate-requests-55gsw-vw92t	0/1	Completed	0	66s
prometheus-example-app-6dcf7dcf8b-hzb19	1/1	Running	0	6h58m
static-app-1-build	0/1	Completed	0	7h38m
static-app-6944df7b7b-4zdn1	1/1	Running	0	12s
static-app-6944df7b7b-6lfpz	1/1	Running	0	12s
static-app-6944df7b7b-752qk	1/1	Running	0	26s
static-app-6944df7b7b-9pgrl	1/1	Running	0	26s
static-app-6944df7b7b-g2s6s	1/1	Running	0	26m
static-app-6944df7b7b-h2qjm	1/1	Running	0	27s
static-app-6944df7b7b-ldpl6	1/1	Running	0	26s
static-app-6944df7b7b-q7qpc	1/1	Running	0	12s
static-app-6944df7b7b-zdd5b	1/1	Running	0	12s

Metrics will now show a spike on the custom metric.

